

A semaphore  $S$  supports two atomic operations:

**$S \rightarrow \text{Wait}()$ :** The process that issues a wait, waits until semaphore  $S$  is available

**$S \rightarrow \text{Signal}()$ :** The process that issues a signal, notifies other processes that  $S$  is free. If processes are waiting, the OS wakes one up.

**A Binary Semaphore** guarantees mutual exclusive access to a resource (only one process enters the critical section at a time). It is usually initialized to 1.

**Too Much Milk:**

Thread A	Thread B
$S \rightarrow \text{Wait}();$	$S \rightarrow \text{Wait}();$
if (noMilk)	if (noMilk)
buy milk;	buy milk;
$S \rightarrow \text{Signal}();$	$S \rightarrow \text{Signal}();$

**A Counting Semaphore** represents a resource with many units available. The initial count to which the semaphore is initialized is usually the number of resources. A counting semaphore lets a process continue as long as more instances are available.

**Semaphores are good for implementing:**

- Scheduling constraints like waitpid or thread join (initial count == 0)
- Mutual exclusion (initial count == 1)
- Multi-instance resources like bounded buffers (initial count > 1)

Problem statement:

- An object is shared among many threads, most only read the object but some write it
- To get good performance we want to allow **multiple readers** at a time
- To get correct operation we want **only one writer** at a time (and zero readers when there is a writer)
- How do we control access to the object to permit this protocol?

## Motivation

---

Reader / writer asymmetry is common!

A few examples:

- Book database at the library
  - Most accesses are searches (read-only)
  - Less often there are checkouts, returns, new books, etc. (writes)
- Making stock market data available on the web
  - Thousands of clients view the data (read-only)
  - A few times per hour the data is updated (writes)
  - If writes are not atomic clients will sometimes get an inconsistent view of the data — potentially an expensive problem
- Render farm
  - Rendering machines need constant, high-bandwidth access to models
  - Models are infrequently updated
- Anonymous CVS
  - Many people view the sources (for `gcc`, for example)
  - Only a few people contribute

## Motivation

---

Readers/writers is purely a performance optimization. Since it adds complexity how can we tell when it is needed?

### Example:

You are running the CVS server for `gcc`, the GNU C Compiler.

- It is important for people to check out a consistent snapshot of the sources
  - `gcc` contains 18,000 files, 177 MB — 22 minutes to sync up with the sources using a fast cable modem
- ⇒ At most 65 readers per day if there is no concurrency between readers
- If there are more readers than this, reader/writer locks are a necessity

**Summary:** Use reader/writer locks when mutual exclusion is too restrictive, and prevents performance goals from being met.

## Reader/Writer Implementation using Semaphores

---

```
class ReadWrite {
public:
    void Read();
    void Write();
private:
    int      readers; // counts readers
    Semaphore my_read_lock; // controls access to readers
    Semaphore wrt;      // toggles entry to first
} // writer or reader

ReadWrite::ReadWrite {
    readers      = 0;
    mutex->count = 1; // mutex
    wrt->count   = 1; // mutex
}

ReadWrite::Write(){
    wrt->Wait(); // any writers or readers?
    <perform write>
    wrt->Signal(); // enable write or read
}

ReadWrite::Read(){
    my_read_lock->Wait(); // reader mutual exclusion
    readers += 1;
    if (readers == 1) // first reader
        wrt->Wait(); // blocks writers
    my_read_lock->Signal();
    <perform read>
    my_read_lock->Wait(); // reader mutual exclusion
    readers -= 1;
    if (readers == 0) //
        wrt->Signal(); // enable writers or readers
    my_read_lock->Signal();
}
```

## Readers/Writers Scenario 1

---

R1:	R2:	W1:
Read ()		
	Read ()	
		Write ()

## Readers/Writers Scenario 2

---

R1:	R2:	W1:
		Write ()
Read ()		
	Read ()	

## Reader/Writers Scenario 3

---

R1:	R2:	W1:
Read ()		
	Read ()	
		Write ()

## Readers/Writers Solution Discussion

---

Implementation notes:

1. The first reader blocks if there is a writer; any other readers who try to enter block on mutex
2. The last reader to exit signals a waiting writer
3. When a writer exits, if there is both a reader and writer waiting, which goes next depends on the scheduler
4. If a writer exits and a reader goes next, then all readers that are waiting will fall through (at least one is waiting on wrt and zero or more can be waiting on mutex)
5. Does this solution guarantee all threads will make progress?

Alternative desirable semantics:

- Let a writer enter its critical section as soon as possible

## Readers/Writers Solution Favoring Writers

---

```
ReadWrite::Write () {
    my_write_lock->Wait(); // writer mutual exclusion
    writers += 1;         // pending writer
    if (writers == 1)     // block readers
        block_readers->Wait();
    my_write_lock->Signal();

    wrt->Wait();          // writer/reader
    <perform write>       // mutual exclusion
    wrt ->Signal();

    my_write_lock->Wait(); // writer mutual exclusion
    writers -= 1;
    if (writers == 0)     // sync with readers
        block_readers->Signal();
    my_write_lock->Signal();
}

ReadWrite::Read () {
    block_readers->Wait(); // block if there's a writer
    my_read_lock->Wait();  // reader mutual exclusion
    readers += 1;         //
    if (readers == 1)     // synchronize with writers
        wrt->Wait();
    my_read_lock->Signal();
    block_readers->Signal();

    <perform read>
    my_read_lock->Wait();  // reader mutual exclusion
    readers -= 1;         // reader done
    if (readers == 0)     // enable writers
        wrt->Signal();
    my_read_lock->Signal();
}
```

### Readers/Writers Scenario 4

---

R1:	R2:	W1:	W2:
Read ( )			
	Read ( )		
		Write ( )	
			Write ( )

### Readers/Writers Scenario 5

---

R1:	R2:	W1:	W2:
		Write ( )	
Read ( )			
	Read ( )		
			Write ( )

## Reader/Writers Scenario 6

---

R1:	R2:	W1:	W2:
Read ()			
		Write ()	
	Read ()		
			Write ()